

Mini Project

Priyanka Bhosale

2024-03-15

Project-3: Using 2 datasets: Home Depot Ad Spending and Google Trends for the analysis.

```
# Load the necessary libraries
library(ggplot2)
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

# Set a CRAN mirror manually
options(repos = structure(c(CRAN = "https://cloud.r-project.org/")))
install.packages('prophet')

##
## The downloaded binary packages are in
## /var/folders/g_/zw8t9hgn0p197fxcrlghncc0000gn/T//RtmpMKVxb8/downloaded_packages
library(prophet)

## Loading required package: Rcpp
## Loading required package: rlang
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

# Read the Google Trends data
google_trends = read.csv("homedepot_googletrends.csv")

# Read the Home Depot ad spending data
homedepot_adspend = read.csv("homedepot.adspend.csv")
```

Original Datasets details:

There are 1821 rows and 5 columns in Google Trends CSV dataset. There are 756 rows and 13 columns in Home Depot Ad Spending dataset.

```
str(google_trends)
```

```
## 'data.frame': 1821 obs. of 5 variables:
## $ X : int 1820 1819 1818 1817 1816 1815 1814 1813 1812 1811 ...
## $ date : chr "Oct 01 2018" "Sep 30 2018" "Sep 29 2018" "Sep 28 2018" ...
## $ value : num 15349 19085 19124 14549 14161 ...
## $ period : chr "2018-10-01" "2018-09-30" "2018-09-29" "2018-09-28" ...
## $ onediffvalue: num NA 3736.4 39.1 -4575 -387.9 ...
```

```
str(homedepot_adspend)
```

```
## 'data.frame': 756 obs. of 13 variables:
## $ TIME.PERIOD : chr "WEEK OF OCT 07, 2013 (B)" "WEEK OF OCT 07, 2013 (B)" "WEEK OF OCT 07, 2013 (B)" ...
## $ PRODUCT : chr "Home Depot Home Center & Kidde : Combo Vignette" "Home Depot Home Center & Kidde : Combo Vignette" ...
## $ TOTAL.DOLS..000. : num 6.3 5128.7 2.6 6.1 4760.4 ...
## $ NETWORK.TV.DOLS..000. : num 0 1465 0 0 1753 ...
## $ CABLE.TV.DOLS..000. : num 0 1211.8 2.6 0 1105.8 ...
## $ SYNDICATION.DOLS..000. : num 0 171.9 0 0 91.1 ...
## $ SPOT.TV.DOLS..000. : num 6.3 348.3 0 6.1 116.9 ...
## $ MAGAZINES.DOLS..000. : num 0 0 0 0 0 0 0 0 0 ...
## $ NATL.NEWSP.DOLS..000. : num 0 0 0 0 0 0 0 0 0 ...
## $ NEWSPAPER.DOLS..000. : num 0 238 0 0 0 ...
## $ NETWORK.RADIO.DOLS..000. : num 0 1635 0 0 1635 ...
## $ NAT.SPOT.RADIO.DOLS..000. : num 0 58.3 0 0 58.3 0 0 58.3 0 0 ...
## $ OUTDOOR.DOLS..000. : num 0 0 0 0 0 0 0 0 0 ...
```

Pre-processing on Google Trends dataset:

There is one missing value, and after dropping this missing row, the dataset now has 1820 rows and 4 columns. The first column is renamed to "id". Column "date" is a character datatype, so converted this into Date datatype. Columns "date" and "period" hold the same data, i.e., the date. So dropped "period" column to avoid having duplicate and irrelevant columns in this dataset. Converted columns "value" and "onediffvalue" to numeric datatypes.

Pre-processing on Home Depot Ad Spending dataset:

There is no missing value. The column names are renamed to make more sense and meaningful, for eg. DOLS renamed to Dollars, removed (000) from the column names. Column TIME_PERIOD is in a format that cannot be used directly for analysis. It is in the format "WEEK OF OCT 07, 2013 (B)". So removed WEEK OF and the last (B) from such values and converted this character datatype to Date datatype with a consistent format.

```
sum(is.na(google_trends))
```

```
## [1] 1
```

```
sum(is.na(homedepot_adspend))
```

```
## [1] 0
```

```

# Get location of null values
which(is.na(google_trends))

## [1] 7285

# Remove rows with missing values
google_trends = na.omit(google_trends)

# Check if the missing data is deleted. This should print 0 now
sum(is.na(google_trends))

## [1] 0

# Add column name to the first column of Google Trends dataset
colnames(google_trends)[1] = "id"

# Rename columns of Home Depot Ad Spending
names(homedepot_adspend) <- c("TIME_PERIOD", "PRODUCT", "TOTAL_DOLLARS", "NETWORK_TV_DOLLARS",
                             "CABLE_TV_DOLLARS", "SYNDICATION_DOLLARS", "SPOT_TV_DOLLARS", "MAGAZINES_DOLLARS",
                             "NATL_NEWSP_DOLLARS", "NEWSPAPER_DOLLARS", "NETWORK_RADIO_DOLLARS",
                             "NAT_SPOT_RADIO_DOLLARS", "OUTDOOR_DOLLARS")

# Check the new column names
colnames(homedepot_adspend)

## [1] "TIME_PERIOD"          "PRODUCT"              "TOTAL_DOLLARS"
## [4] "NETWORK_TV_DOLLARS"   "CABLE_TV_DOLLARS"     "SYNDICATION_DOLLARS"
## [7] "SPOT_TV_DOLLARS"      "MAGAZINES_DOLLARS"    "NATL_NEWSP_DOLLARS"
## [10] "NEWSPAPER_DOLLARS"    "NETWORK_RADIO_DOLLARS" "NAT_SPOT_RADIO_DOLLARS"
## [13] "OUTDOOR_DOLLARS"

# Convert date column to Date format
google_trends$date = as.Date(google_trends$date, format = "%b %d %Y")

# Remove label "WEEK OF" and the last "(B)" from the "TIME PERIOD" column of
# Home Depot Ad Spending dataset
homedepot_adspend$`TIME_PERIOD` = trimws(sub("WEEK OF ", "", homedepot_adspend$`TIME_PERIOD`))
homedepot_adspend$`TIME_PERIOD` = trimws(sub(" \\(B\\)$", "", homedepot_adspend$`TIME_PERIOD`))
homedepot_adspend$`TIME_PERIOD` = as.Date(homedepot_adspend$`TIME_PERIOD`, format = "%b %d, %Y")

# Since we converted date column to Date datatype, "period" column is a duplicate and
# irrelevant column. Drop "period".
drop = c('period')
google_trends = google_trends[,!(names(google_trends) %in% drop)]

# Convert the value and onediffvalue columns to numeric
google_trends$value <- as.numeric(google_trends$value)
google_trends$onediffvalue <- as.numeric(google_trends$onediffvalue)

str(google_trends)

## 'data.frame': 1820 obs. of 4 variables:
## $ id : int 1819 1818 1817 1816 1815 1814 1813 1812 1811 1810 ...
## $ date : Date, format: "2018-09-30" "2018-09-29" ...
## $ value : num 19085 19124 14549 14161 15195 ...
## $ onediffvalue: num 3736.4 39.1 -4575 -387.9 1033.2 ...

```

```
str(homedepot_adspend)
```

```
## 'data.frame': 756 obs. of 13 variables:
## $ TIME_PERIOD : Date, format: "2013-10-07" "2013-10-07" ...
## $ PRODUCT : chr "Home Depot Home Center & Kidde : Combo Vignette" "Home Depot Home C
## $ TOTAL_DOLLARS : num 6.3 5128.7 2.6 6.1 4760.4 ...
## $ NETWORK_TV_DOLLARS : num 0 1465 0 0 1753 ...
## $ CABLE_TV_DOLLARS : num 0 1211.8 2.6 0 1105.8 ...
## $ SYNDICATION_DOLLARS : num 0 171.9 0 0 91.1 ...
## $ SPOT_TV_DOLLARS : num 6.3 348.3 0 6.1 116.9 ...
## $ MAGAZINES_DOLLARS : num 0 0 0 0 0 0 0 0 0 ...
## $ NATL_NEWSP_DOLLARS : num 0 0 0 0 0 0 0 0 0 ...
## $ NEWSPAPER_DOLLARS : num 0 238 0 0 0 ...
## $ NETWORK_RADIO_DOLLARS : num 0 1635 0 0 1635 ...
## $ NAT_SPOT_RADIO_DOLLARS : num 0 58.3 0 0 58.3 0 0 58.3 0 0 ...
## $ OUTDOOR_DOLLARS : num 0 0 0 0 0 0 0 0 0 ...
```

Summarize Google Trends Data:

```
summary(google_trends)
```


##	id	date	value	onediffvalue
##	Min. :	0.0	Min. :2013-10-07	Min. : 9205
##	1st Qu.:	454.8	1st Qu.:2015-01-04	Min. :-10129.249
##	Median :	909.5	Median :2016-04-03	1st Qu.: -892.936
##	Mean :	909.5	Mean :2016-04-03	Median : 21.625
##	3rd Qu.:	1364.2	3rd Qu.:2017-07-02	Mean : -1.292
##	Max. :	1819.0	Max. :2018-09-30	3rd Qu.: 853.101
##				Max. : 13500.754

Visualize Google Trends Data:

This plot shows the Google Trends data, specifically the trend in the search volume for “Home Depot” over time. The x-axis represents the date, and the y-axis represents the value. The plot shows a clear seasonal pattern, with value peaking around the same time each year, likely due to seasonal factors affecting home improvement projects.

```
# Visualize the Google Trends data -> Date vs Value
ggplot(data = google_trends, aes(x = date, y = value)) +
  geom_line() +
  labs(title = "Google Trends Data Visualization")
```

Google Trends Data Visualization



```
# Performing adf test for checking the stationarity  
library(zoo)
```

```
##  
## Attaching package: 'zoo'  
## The following objects are masked from 'package:base':  
##  
##   as.Date, as.Date.numeric
```

```
library(tseries)
```

```
# Perform the ADF test  
adf_result <- adf.test(google_trends$value)
```

```
## Warning in adf.test(google_trends$value): p-value smaller than printed p-value
```

```
# Print the ADF test result  
print(adf_result)
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: google_trends$value  
## Dickey-Fuller = -5.3037, Lag order = 12, p-value = 0.01  
## alternative hypothesis: stationary
```

Even though p-value is less than

α

(where


α

=0.05), we can see a significant variance in data. So to get rid off the variance we need to difference the data. There is a differenced data in google trends dataset and is stored in "onediffvalue" column. We now will perform plot the graph for onediffvalue.

Visualize Google Trends Data -> Date vs One Difference Value:

This plot shows the first-order differenced values of the Google Trends data, which helps to remove the trend and make the time series stationary. The x-axis represents the date, and the y-axis represents the differenced value. The plot shows a more stable pattern around a constant mean, indicating that the differencing has removed the trend and seasonality from the original data.

```
# Line plot for x = date and y = onediffvalue
ggplot(data = google_trends, aes(x = date, y = onediffvalue)) +
  geom_line() +
  labs(title = "Google Trends: Date vs. One Diff Value", x = "Date", y = "One Diff Value")
```




ACF and PACF Plots for Google Trends with One Difference Value:

These plots show the Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) for the differenced Google Trends data. The ACF plot shows a significant spike at lag 1, suggesting the presence of a first-order moving average (MA) component in the time series. The PACF plot shows a decaying pattern, indicating the potential presence of an autoregressive (AR) component. These plots help identify the appropriate orders for the ARIMA/SARIMA models.

```
# ACF and PACF plots for Google Trends
```


```
acf(google_trends$onediffvalue, main = "ACF Plot for Google Trends with One Difference Value")
```

ACF Plot for Google Trends with One Difference Value



```
pacf(google_trends$onediffvalue, main = "PACF Plot for Google Trends with One Difference Value")
```

PACF Plot for Google Trends with One Difference Value



Summarize Home Depot Advertisement Spending Data:

```
summary(homedepot_adspend)
```

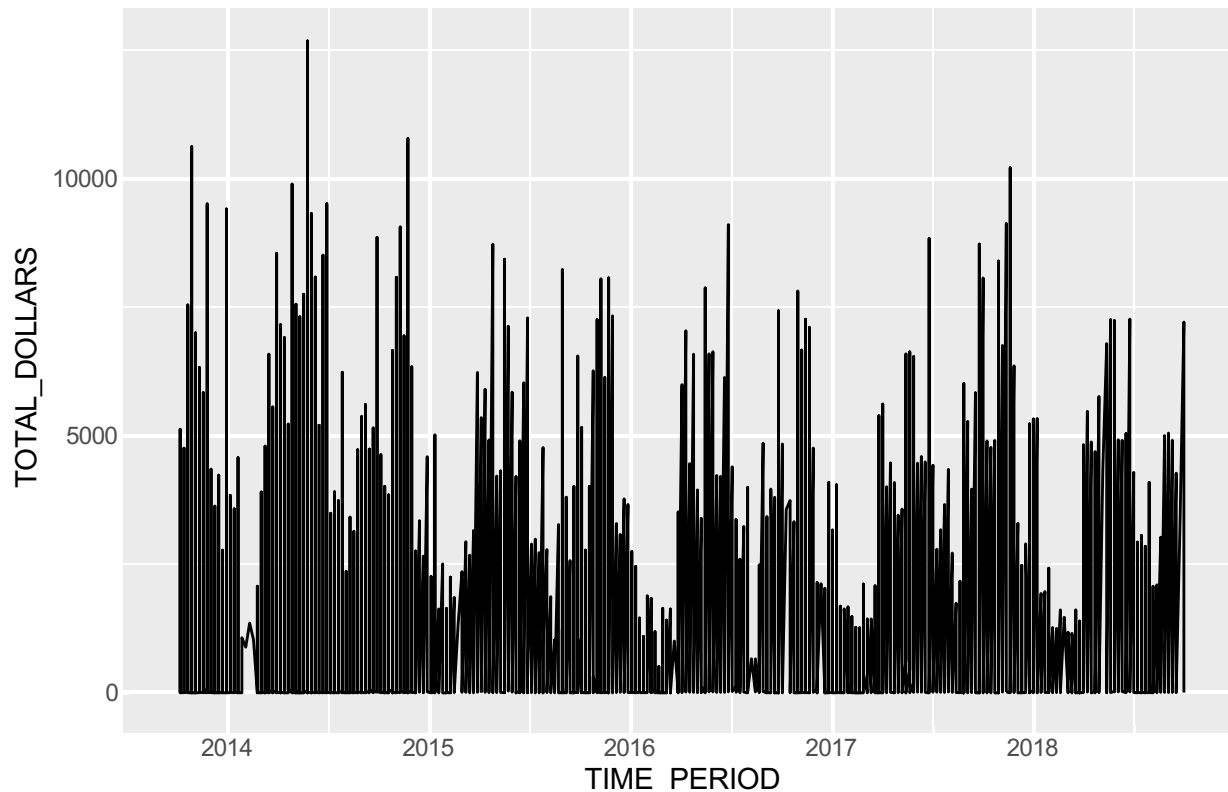
```
##  TIME_PERIOD      PRODUCT      TOTAL_DOLLARS  NETWORK_TV_DOLLARS
##  Min.   :2013-10-07  Length:756      Min.    :   0.0    Min.    :   0
##  1st Qu.:2014-10-20   Class :character  1st Qu.:   4.1    1st Qu.:   0
##  Median :2015-12-28   Mode  :character  Median :  13.1    Median :   0
##  Mean   :2016-02-06                      Mean   :1583.4    Mean   : 398
##  3rd Qu.:2017-05-02                      3rd Qu.:2862.2    3rd Qu.:   2
##  Max.   :2018-10-01                      Max.   :12679.8   Max.   :6558
##  CABLE_TV_DOLLARS SYNDICATION_DOLLARS SPOT_TV_DOLLARS  MAGAZINES_DOLLARS
##  Min.    :   0.0    Min.    :   0.0    Min.    :   0.00   Min.    :   0.0
##  1st Qu.:   0.0    1st Qu.:   0.0    1st Qu.:   0.00   1st Qu.:   0.0
##  Median :  10.3    Median :   0.0    Median :   0.30   Median :   0.0
##  Mean   : 552.9    Mean   :  12.2    Mean   :  27.76   Mean   : 125.9
##  3rd Qu.:1058.5    3rd Qu.:  10.5    3rd Qu.:  23.07   3rd Qu.:   0.0
##  Max.   :4208.1    Max.   :  418.3    Max.   :849.70   Max.   :4440.9
##  NATL_NEWSP_DOLLARS NEWSPAPER_DOLLARS NETWORK_RADIO_DOLLARS
##  Min.    : 0.0000    Min.    :   0.00   Min.    :   0.0
##  1st Qu.: 0.0000    1st Qu.:   0.00   1st Qu.:   0.0
##  Median : 0.0000    Median :   0.00   Median :   0.0
##  Mean   : 0.3985    Mean   :  40.45   Mean   : 372.1
##  3rd Qu.: 0.0000    3rd Qu.:   2.55   3rd Qu.: 784.2
##  Max.   :61.3000    Max.   :958.40   Max.   :2485.6
##  NAT_SPOT_RADIO_DOLLARS OUTDOOR_DOLLARS
##  Min.    :   0.00    Min.    :   0.000
##  1st Qu.:   0.00    1st Qu.:   0.000
##  Median :   0.00    Median :   0.000
##  Mean   : 51.57     Mean   :   2.043
##  3rd Qu.: 18.80     3rd Qu.:   0.000
##  Max.   :848.40     Max.   :226.300
```

Visualize Home Depot Advertisement Spending Data -> Time Period vs Total Dollars:

This plot shows the Home Depot advertisement spending over time, represented by the Total Dollars spent on advertising. The x-axis represents the time period, and the y-axis represents the total advertising dollars spent. The plot reveals a cyclical pattern, with peaks and valleys in advertising spending, possibly related to seasonal demand or marketing campaigns.

```
# Visualize the Home Depot Ad Spending data
ggplot(data = homedepot_adspend, aes(x = TIME_PERIOD, y = TOTAL_DOLLARS)) +
  geom_line() +
  labs(title = "Home Depot Ad Spending Data Visualization")
```


Home Depot Ad Spending Data Visualization



ADF test for Home Depot Ad Spending TOTAL_DOLLARS:

The plot and the ADF test signifies stationarity in the data. We go ahead and take log and differencing to remove/minimize stationarity.

```
# Performing adf test for checking the stationarity
```

```
adf_result <- adf.test(homedepot_adspend$TOTAL_DOLLARS)
```

```
## Warning in adf.test(homedepot_adspend$TOTAL_DOLLARS): p-value smaller than  
## printed p-value
```

```
# Print the ADF test result
```

```
print(adf_result)
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

```
## data: homedepot_adspend$TOTAL_DOLLARS
```

```
## Dickey-Fuller = -5.8114, Lag order = 9, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

First-order differencing for TOTAL_DOLLARS:

This plot shows the first-order differenced values of the Home Depot advertisement spending, which helps to remove the trend and make the time series stationary. The x-axis represents the time period, and the y-axis represents the differenced total advertising dollars. The plot shows a more stable pattern around a constant

mean, indicating that the differencing has removed the trend and seasonality from the original data.


```
homedepot_adspend$DIFF_TOTAL_DOLLARS = c(NA, diff(log(homedepot_adspend$TOTAL_DOLLARS)))
```

```
# Visualize the Home Depot Ad Spending data
```

```
ggplot(data = homedepot_adspend, aes(x = TIME_PERIOD, y = DIFF_TOTAL_DOLLARS)) +  
  geom_line() +  
  labs(title = "Differenced Home Depot Ad Spending Data Visualization")
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```

Differenced Home Depot Ad Spending Data Visualization



ACF and PACF with DIFF_TOTAL_DOLLARS:

These plots show the ACF and PACF for the differenced Home Depot advertisement spending data. The ACF plot shows a significant spike at lag 1, suggesting the presence of a first-order moving average (MA) component in the time series. The PACF plot shows a decaying pattern, indicating the potential presence of an autoregressive (AR) component. These plots help identify the appropriate orders for the ARIMA/SARIMA models for the advertisement spending data.

```
sum(is.na(homedepot_adspend))
```

```
## [1] 1
```

```
homedepot_adspend = na.omit(homedepot_adspend)
```


```
# Filter rows with infinite values in DIFF_TOTAL_DOLLARS
```

```
inf_rows <- which(!is.finite(homedepot_adspend$DIFF_TOTAL_DOLLARS))
```

```
# Remove rows with infinite values
homedepot_adspend <- homedepot_adspend[-inf_rows, ]


# ACF and PACF plots for Google Trends
acf(homedepot_adspend$DIFF_TOTAL_DOLLARS, main = "ACF Plot for Homedepot's DIFF_TOTAL_DOLLARS")
```

ACF Plot for Homedepot's DIFF_TOTAL_DOLLARS



```
pacf(homedepot_adspend$DIFF_TOTAL_DOLLARS, main = "PACF Plot for Homedepot's DIFF_TOTAL_DOLLARS")
```

PACF Plot for Homedepot's DIFF_TOTAL_DOLLARS



Relationship between Google Trends and advertisement spending:

Merge the 2 datasets into one, on columns "date" of Google Trends and "Time_Period" of Home Depot Ads Spending. Removed missing values.

Google Trends's "date" column has daily data and Home Depot Ads Spending's "TIME_PERIOD". So when these 2 datasets are merged on date and TIME_PERIOD columns, we get the merged dataset with weekly data with 735 rows. Head() of this merged data is also printed for reference that shows date as weekly.

```
# Merge the datasets based on the date column
merged_data <- merge(google_trends, homedepot_adspend, by.x = "date", by.y = "TIME_PERIOD", all = TRUE)
sum(is.na(merged_data))
```

```
## [1] 20286
```

```
# Remove rows with missing values
merged_data = na.omit(merged_data)
```

```
print(nrow(merged_data))
```

```
## [1] 735
```

```
head(merged_data)
```

```
##      date id  value onediffvalue
## 1 2013-10-07 0 12997.16      222.3070
## 2 2013-10-07 0 12997.16      222.3070
## 9 2013-10-14 7 13957.23     1377.4804
## 10 2013-10-14 7 13957.23     1377.4804
## 11 2013-10-14 7 13957.23     1377.4804
## 18 2013-10-21 14 13282.89      733.4774
```


##	PRODUCT	TOTAL_DOLLARS
## 1	Home Depot Home Center : Home Center/Hardware Store	5128.7
## 2	Home Depot Home Center : Vignette	2.6
## 9	Home Depot Home Center : Vignette	13.9
## 10	Home Depot Home Center : Home Center/Hardware Store	4760.4
## 11	Home Depot Home Center & Kidde : Combo Vignette	6.1
## 18	Home Depot Home Center : Vignette	4.4
##	NETWORK_TV_DOLLARS CABLE_TV_DOLLARS SYNDICATION_DOLLARS SPOT_TV_DOLLARS	
## 1	1465.2 1211.8 171.9 348.3	
## 2	0.0 2.6 0.0 0.0	
## 9	0.0 13.9 0.0 0.0	
## 10	1753.1 1105.8 91.1 116.9	
## 11	0.0 0.0 0.0 6.1	
## 18	0.0 3.9 0.0 0.5	
##	MAGAZINES_DOLLARS NATL_NEWSP_DOLLARS NEWSPAPER_DOLLARS NETWORK_RADIO_DOLLARS	
## 1	0 0 238.2 1635.2	
## 2	0 0 0.0 0.0	
## 9	0 0 0.0 0.0	
## 10	0 0 0.0 1635.2	
## 11	0 0 0.0 0.0	
## 18	0 0 0.0 0.0	
##	NAT_SPOT_RADIO_DOLLARS OUTDOOR_DOLLARS DIFF_TOTAL_DOLLARS	
## 1	58.3 0 6.7020579	
## 2	0.0 0 -7.5870960	
## 9	0.0 0 -5.8361981	
## 10	58.3 0 6.6597982	
## 11	0.0 0 0.8527773	
## 18	0.0 0 -7.4470756	

Visualize Merged Data:


1. Plot-1 (Differenced Google Trends Differenced Value vs. Date): This plot shows the differenced Google Trends data over time. The x-axis represents the date, and the y-axis represents the differenced value. This plot is similar to the "Google Trends: Date vs. One Diff Value" plot but is presented after merging the Google Trends and Home Depot advertisement spending datasets.
2. Plot-2 (Differenced TOTAL DOLLARS for Home Depot Ad Spending Data): This plot shows the differenced Home Depot advertisement spending over time. The x-axis represents the date, and the y-axis represents the differenced total advertising dollars. This plot is similar to the "Differenced Home Depot Ad Spending Data Visualization" but is presented after merging the Google Trends and Home Depot advertisement spending datasets.
3. Plot-3 (Google Trends Differenced Value vs. Total Advertisement Spending Differenced Total Dollars): This plot shows the relationship between the differenced Google Trends data (y-axis) and the differenced Home Depot advertisement spending (x-axis). It helps visualize the potential correlation or relationship between the two variables, which is useful for building regression models like SARIMAX.

```
# Plot the differenced Google Trends data
ggplot(merged_data, aes(x = date, y = onediffvalue)) +
  geom_line() +
  labs(title = "Differenced Google Trends Differenced Value vs Date", x = "Date",
        y = "Differenced Value")
```

Differenced Google Trends Differenced Value vs Date

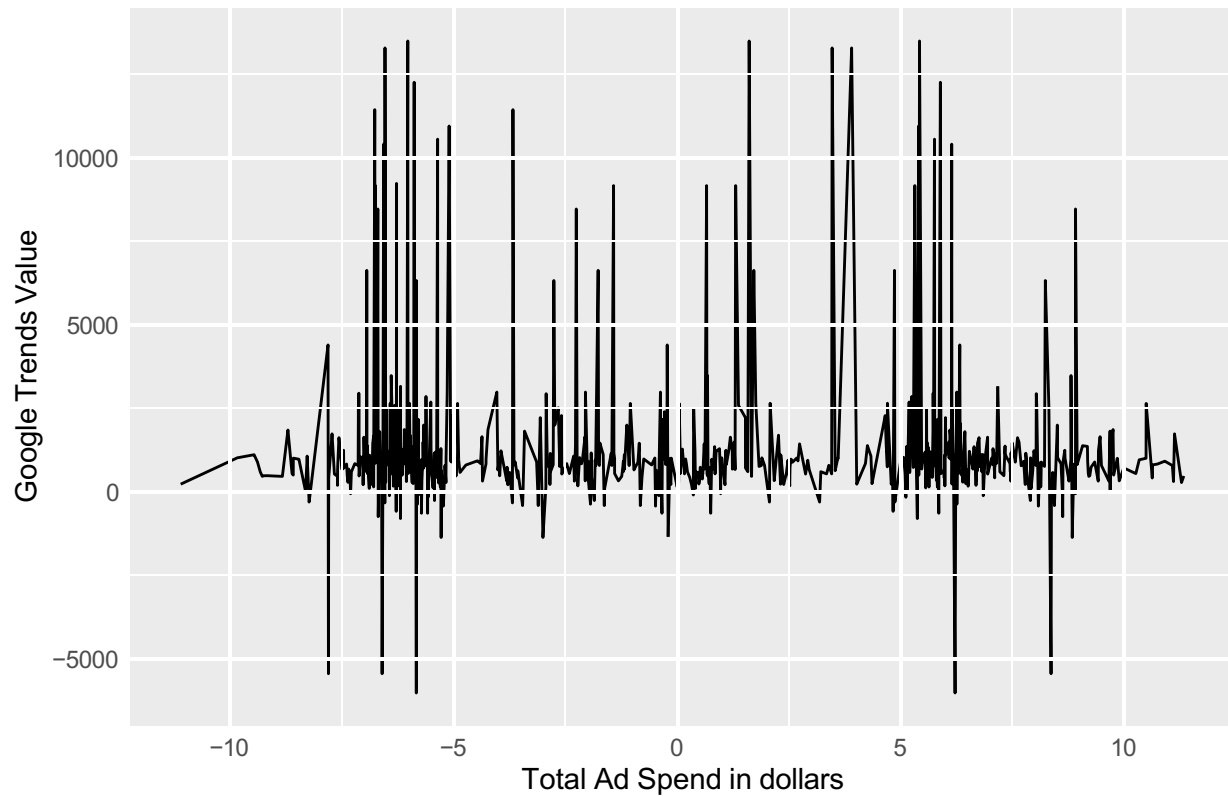


```
# Plotting the differenced Home Depot Ad Spending data
ggplot(merged_data, aes(x = date, y = DIFF_TOTAL_DOLLARS)) +
  geom_line() +
  labs(title = "Differenced TOTAL DOLLARS for Home Depot Ad Spending Data", x = "Date",
        y = "Differenced Dollars")
```



```
# Plot Google Trends data over time
ggplot(data = merged_data, aes(x = DIFF_TOTAL_DOLLARS, y = onediffvalue)) +
  geom_line() +
  labs(title = "Google Trends Differenced Value vs Total Advertisement Spending Differenced Total Dolla
           x = "Total Ad Spend in dollars", y = "Google Trends Value")
```


Google Trends Differenced Value vs Total Advertisement Spending Differ



ACF and PACF plots for the merged dataset - Differenced values of Google Trends (onediffvalue) and Home Depot Ad Spending (DIFF_TOTAL_DOLLARS)


```
acf(merged_data$onediffvalue, main = "ACF Plot for merged dataset: onediffvalue")
```


ACF Plot for merged dataset: onediffvalue




```
pacf(merged_data$onediffvalue, main = "PACF Plot for merged dataset: Google Trends's onediffvalue")
```

PACF Plot for merged dataset: Google Trends's onediffvalue




```
acf(merged_data$DIFF_TOTAL_DOLLARS, main = "ACF Plot for merged dataset: DIFF_TOTAL_DOLLARS")
```

ACF Plot for merged dataset: DIFF_TOTAL_DOLLARS



```
pacf(merged_data$DIFF_TOTAL_DOLLARS, main = "PACF Plot for merged dataset: DIFF_TOTAL_DOLLARS")
```

PACF Plot for merged dataset: DIFF_TOTAL_DOLLARS



Split the data into Training and Test (80% and 20%)

```
# Split the data into training and testing sets
train_size <- floor(0.8 * nrow(merged_data))
train_data <- merged_data[1:train_size, ]
test_data <- merged_data[(train_size + 1):nrow(merged_data), ]
```

Train models on train dataset:

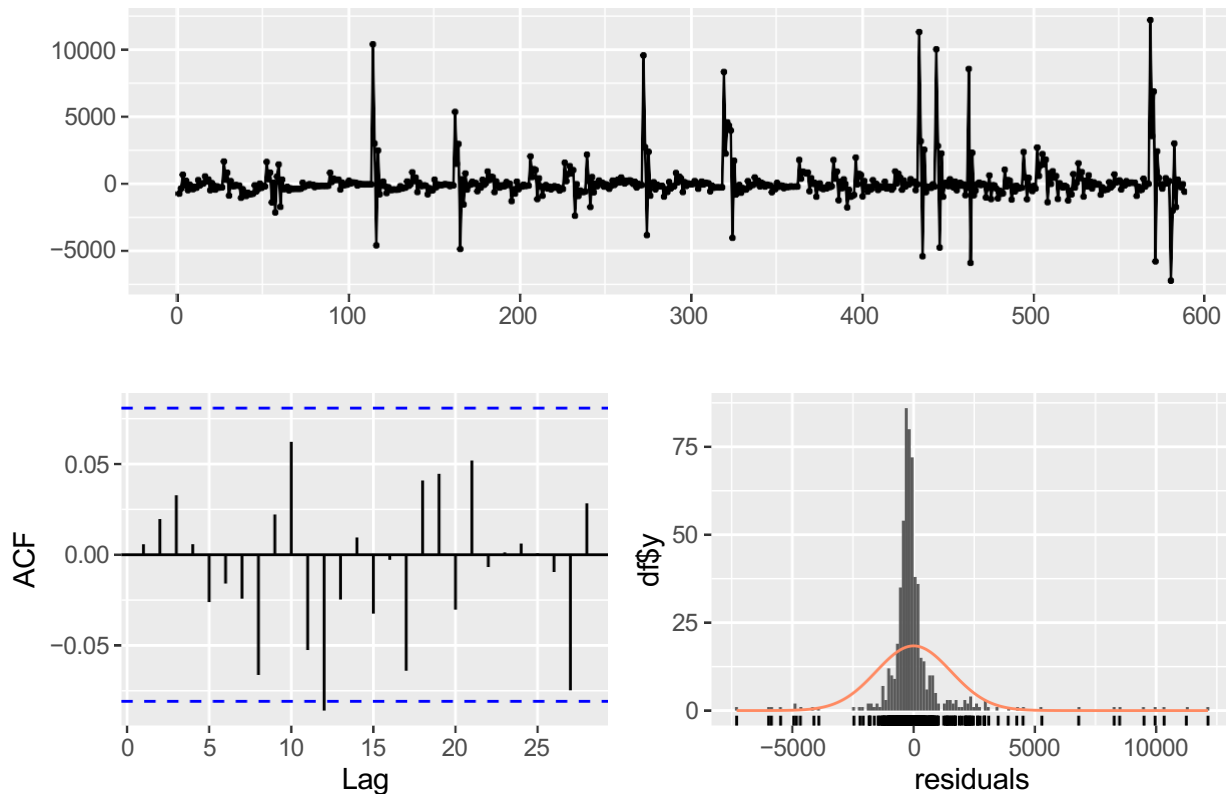
We chose to fit 2 models:

1. SARIMA without additional variables: We use `auto.arima()` to train the model. It correctly identifies the model to be a (0,0,2) model that matches well with the ACF and PACF plots mentioned above for the merged dataset's `onediffvalue`.
2. SARIMAX with additional variables: We are using `DIFF_TOTAL_DOLLARS` of the merged dataset to be our external regressor and `onediffvalue` to be our target variable. The `auto.arima()` chooses the model to be a (0,0,2) model.

```
# Train SARIMA model without additional variables
sarima_model_wo_additional <- auto.arima(train_data$onediffvalue)
summary(sarima_model_wo_additional)

## Series: train_data$onediffvalue
## ARIMA(0,0,2) with non-zero mean
##
## Coefficients:
##          ma1          ma2          mean
##          0.7082    0.2296   1229.8776
## s.e.    0.0400    0.0392    122.7897
##
## sigma^2 = 2377761:  log likelihood = -5149.5
## AIC=10307  AICc=10307.07  BIC=10324.5
##
## Training set error measures:
##              ME    RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 0.3020978 1538.06 702.5965 -82.11875 176.9225 1.347009 0.005785303
checkresiduals(sarima_model_wo_additional)
```

Residuals from ARIMA(0,0,2) with non-zero mean



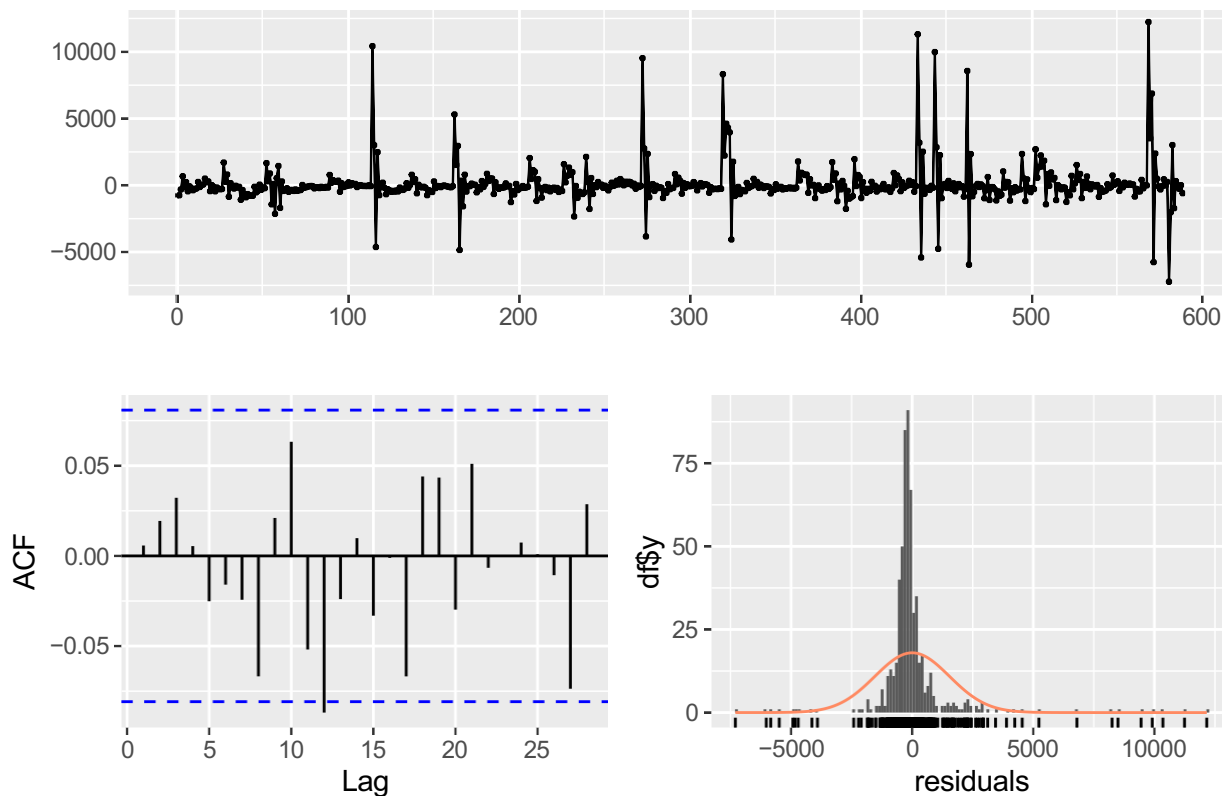
```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,0,2) with non-zero mean
## Q* = 7.0604, df = 8, p-value = 0.5301
##
## Model df: 2.    Total lags used: 10
## Train SARIMAX model with additional variables
sarimax_model_additional <- auto.arima(train_data$onediffvalue,
                                         xreg = train_data$DIFF_TOTAL_DOLLARS, seasonal = TRUE)
summary(sarimax_model_additional)

## Series: train_data$onediffvalue
## Regression with ARIMA(0,0,2) errors
##
## Coefficients:
##      ma1    ma2 intercept    xreg
##    0.7091 0.2307 1229.7295  3.7555
## s.e. 0.0400 0.0393 122.8847  7.5876
##
## sigma^2 = 2380836: log likelihood = -5149.38
## AIC=10308.75 AICc=10308.86 BIC=10330.64
##
## Training set error measures:
##              ME    RMSE    MAE    MPE    MAPE    MASE
## Training set 0.4637945 1537.738 703.3406 -83.77338 178.2107 1.348436
##              ACF1
```

```
## Training set 0.005749775
```

```
checkresiduals(sarimax_model_additional)
```

Residuals from Regression with ARIMA(0,0,2) errors



```
##  
## Ljung-Box test  
##  
## data: Residuals from Regression with ARIMA(0,0,2) errors  
## Q* = 7.0919, df = 8, p-value = 0.5267  
##  
## Model df: 2. Total lags used: 10
```

Evaluate models on test dataset

```
# Evaluate SARIMA model without additional variables  
sarima_model_wo_additional_pred <- forecast(sarima_model_wo_additional, h = nrow(test_data))  
  
# Evaluate SARIMAX model with additional variables  
sarimax_model_additional_pred <- forecast(sarimax_model_additional, h = nrow(test_data),  
                                           xreg = test_data$DIFF_TOTAL_DOLLARS)
```

MSPE for all models:

The MSPE values are calculated for the SARIMA and SARIMAX models on the test dataset, providing a measure of the models' predictive performance.

MSPE for SARIMA model without additional variable: 7558075.43500495

MSPE for SARIMAX model with additional variable: 7556795.88658626

A lower MSPE value indicates better predictive accuracy. There is not much of a difference in this case, but “SARIMAX model with additional variable” performs slightly better than “SARIMA model without additional variable”.

```
# MSPE for SARIMA model without additional variables
```

```
sarima_model_wo_additional_mspe <- mean((sarima_model_wo_additional_pred$mean - test_data$onediffvalue)  
print(paste("Mean Squared Prediction Error (MSPE) for SARIMA model without additional variable:",  
            sarima_model_wo_additional_mspe))
```

```
## [1] "Mean Squared Prediction Error (MSPE) for SARIMA model without additional variable: 7558075.4350
```

```
# MSPE for SARIMAX model with additional variables
```

```
sarimax_model_additional_mspe <- mean((sarimax_model_additional_pred$mean - test_data$onediffvalue)^2)  
print(paste("Mean Squared Prediction Error (MSPE) for SARIMAX model with additional variable:",  
            sarimax_model_additional_mspe))
```

```
## [1] "Mean Squared Prediction Error (MSPE) for SARIMAX model with additional variable: 7556795.886586
```

Facebook Prophet on Google Trends:

Now, we implement the Facebook Prophet model, which is an additive regression model for time series forecasting. The Prophet model is fit to the original Google Trends data, incorporating weekly seasonality. The MSPE value for the Prophet model is calculated on the test dataset, allowing for a comparison with the SARIMA and SARIMAX models. MSPE value for the Prophet model is 6611852.07536799.

```
# Rename columns
```

```
prophet_data = merged_data[, c("date", "value")]  
names(prophet_data) = c('ds', 'y')
```

```
# Make sure 'ds' is a date type
```

```
prophet_data$ds <- as.Date(prophet_data$ds)
```

```
# Fit Prophet model
```

```
prophet_model <- prophet(prophet_data, weekly.seasonality = TRUE, daily.seasonality = FALSE)
```

```
# Prepare future dataframe for predictions
```

```
future <- make_future_dataframe(prophet_model, periods = nrow(test_data))
```

```
# Forecast
```

```
prophet_forecast = predict(prophet_model, future)
```

```
# Extract the predicted values for the test set dates
```

```
forecasted_test_values <- prophet_forecast %>%  
  filter(ds > max(train_data$ds)) %>%  
  select(ds, yhat)
```

```
## Warning: There was 1 warning in `filter()`.  
## i In argument: `ds > max(train_data$ds)`.
```

```
## Caused by warning in `max()`:  
## ! no non-missing arguments to max; returning -Inf
```

```
# Merge the forecasted values with the actual values for the test set
```

```
test_data <- merge(test_data, forecasted_test_values, by.x = "date", by.y = "ds")
```

```
# Calculate MSPE on the test set
```

```
prophet_mspe <- mean((test_data$value - test_data$yhat)^2, na.rm = TRUE)
print(paste("Mean Squared Prediction Error (MSPE) for Prophet:", prophet_mspe))
```

```
## [1] "Mean Squared Prediction Error (MSPE) for Prophet: 6611852.07536799"
```

Compare MSPE's for all 3 models:

Prophet model better than SARIMA and SARIMAX since it has the lowest MSPE as compared to the other 2 models.

```
# Create a data frame with the model names and MSPE values
models <- c("SARIMA", "SARIMAX", "Prophet")
mspe <- c(sarima_model_wo_additional_mspe, sarimax_model_additional_mspe, prophet_mspe)
data <- data.frame(models, mspe)

# Create a bar plot
ggplot(data, aes(x = models, y = mspe, fill = models)) +
  geom_bar(stat = "identity") +
  labs(title = "Comparison of MSPE Values",
       x = "Model",
       y = "MSPE") +
  theme_minimal()
```

